

Product Requirements Document

Title: AI-driven Cloud Cost Copilot (FinOps Assistant)

Problem Statement:

Cloud costs are rising rapidly and becoming unpredictable. Finance and engineering teams struggle to understand cost trends, identify waste, and act quickly to optimize spending.

Target Users:

- **FinOps team / Cloud cost analysts** - analyze trends, forecast, control budget
- **Engineering team leads** - monitor cost spikes, fix underutilized resources
- **Management / Business stakeholders** - get summarized cost insights

Top Use Cases:

1. View monthly cost breakdown and trends per service, account, or team
2. Get automated recommendations to reduce idle or underutilized resources
3. Ask natural-language questions like “Why did compute cost increase in May?”
4. Track tagging/ownership coverage of resources
5. Export KPI reports (CSV/Excel) for monthly review

Success Metrics:

- 90%+ accuracy on KPI calculations (validated against manual SQL)
- Answer 80%+ of benchmark natural-language cost questions correctly
- Detect at least 70% of known cost anomalies in test data
- System runs end-to-end (ETL→ vector store → RAG → API → U locally)
- 1-click run docker-compose up

